

In [1]: *# Import the numpy and pandas package*

```
import numpy as np
import pandas as pd

# Data Visualisation
import matplotlib.pyplot as plt
import seaborn as sns
```

In [8]: `advertising = pd.read_csv("advertising.csv")`
`advertising.head()`

Out[8]:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9

In [9]: `advertising.shape`

Out[9]: (200, 4)

In [10]: `advertising.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0    TV          200 non-null    float64
1    Radio       200 non-null    float64
2    Newspaper   200 non-null    float64
3    Sales       200 non-null    float64
dtypes: float64(4)
memory usage: 6.4 KB
```

In [11]: `advertising.describe()`


Out[11]:

	TV	Radio	Newspaper	Sales
count	200.000000	200.000000	200.000000	200.000000
mean	147.042500	23.264000	30.554000	15.130500
std	85.854236	14.846809	21.778621	5.283892
min	0.700000	0.000000	0.300000	1.600000
25%	74.375000	9.975000	12.750000	11.000000
50%	149.750000	22.900000	25.750000	16.000000
75%	218.825000	36.525000	45.100000	19.050000
max	296.400000	49.600000	114.000000	27.000000


```
In [12]: advertising.isnull().sum()*100/advertising.shape[0]
```

```
Out[12]: TV          0.0  
Radio        0.0  
Newspaper    0.0  
Sales        0.0  
dtype: float64
```


```
In [13]: fig, axs = plt.subplots(3, figsize = (5,5))  
plt1 = sns.boxplot(advertising['TV'], ax = axs[0])  
plt2 = sns.boxplot(advertising['Newspaper'], ax = axs[1])  
plt3 = sns.boxplot(advertising['Radio'], ax = axs[2])  
plt.tight_layout()
```



```
In [14]: sns.boxplot(advertising['Sales'])  
plt.show()
```



```
In [15]: sns.pairplot(advertising, x_vars=['TV', 'Newspaper', 'Radio'], y_vars='Sales', height=10,
plt.show())
```



```
In [17]: X = advertising['TV']
y = advertising['Sales']
```

```
In [18]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size = 0.7, test_size = 0.3)
```

```
In [19]: X_train.head()
```

```
Out[19]: 74      213.4
3        151.5
185      205.0
26       142.9
90       134.3
Name: TV, dtype: float64
```

```
In [20]: y_train.head()
```

```
Out[20]: 74      17.0
          3       16.5
          185     22.6
          26      15.0
          90      14.0
          Name: Sales, dtype: float64
```

```
In [21]: import statsmodels.api as sm
```

```
In [22]: # Add a constant to get an intercept
X_train_sm = sm.add_constant(X_train)

# Fit the regression line using 'OLS'
lr = sm.OLS(y_train, X_train_sm).fit()
```

```
In [23]: lr.params
```

```
Out[23]: const      6.948683
          TV       0.054546
          dtype: float64
```

```
In [24]: print(lr.summary())
```


```

                        OLS Regression Results
=====
Dep. Variable:          Sales    R-squared:                0.816
Model:                  OLS      Adj. R-squared:            0.814
Method:                 Least Squares    F-statistic:          611.2
Date:                  Tue, 09 Apr 2024    Prob (F-statistic):    1.52e-52
Time:                  21:44:45    Log-Likelihood:        -321.12
No. Observations:      140    AIC:                   646.2
Df Residuals:          138    BIC:                   652.1
Df Model:               1
Covariance Type:       nonrobust
=====
               coef      std err          t      P>|t|      [0.025      0.975]
-----
const          6.9487      0.385     18.068      0.000        6.188        7.709
TV              0.0545      0.002     24.722      0.000        0.050        0.059
=====
Omnibus:              0.027    Durbin-Watson:           2.196
Prob(Omnibus):         0.987    Jarque-Bera (JB):         0.150
Skew:                 -0.006    Prob(JB):                 0.928
Kurtosis:              2.840    Cond. No.                  328.
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [25]: plt.scatter(X_train, y_train)
          plt.plot(X_train, 6.948 + 0.054*X_train, 'r')
          plt.show()
```



```
In [26]: y_train_pred = lr.predict(X_train_sm)
         res = (y_train - y_train_pred)
```

```
In [27]: fig = plt.figure()
         sns.distplot(res, bins = 15)
         fig.suptitle('Error Terms', fontsize = 15)           # Plot heading
         plt.xlabel('y_train - y_train_pred', fontsize = 15)  # X-Label
         plt.show()
```

C:\Users\Admin\AppData\Local\Temp\ipykernel_9992\3003513444.py:2: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(res, bins = 15)
```


C:\Users\Admin\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

```
with pd.option_context('mode.use_inf_as_na', True):
```

Error Terms



```
In [28]: plt.scatter(X_train,res)
plt.show()
```



```
In [29]: # Add a constant to X_test
X_test_sm = sm.add_constant(X_test)

# Predict the y values corresponding to X_test_sm
y_pred = lr.predict(X_test_sm)
y_pred.head()
```

```
Out[29]: 126    7.374140
        104    19.941482
        99    14.323269
        92    18.823294
        111   20.132392
        dtype: float64
```

```
In [30]: from sklearn.metrics import mean_squared_error
        from sklearn.metrics import r2_score
```

```
In [31]: #Returns the mean squared error; we'll take a square root
np.sqrt(mean_squared_error(y_test, y_pred))
```

```
Out[31]: 2.019296008966232
```


```
In [32]: r_squared = r2_score(y_test, y_pred)
        r_squared
```

```
Out[32]: 0.7921031601245659
```

```
In [33]: plt.scatter(X_test, y_test)
        plt.plot(X_test, 6.948 + 0.054 * X_test, 'r')
```



```
plt.show()
```



In []: